

```
# Exploration
# An agent primarily improves its knowledge about each activity by gathering more information to make best overall actions.

# Exploitation
# An agent aims at more rewards based on the estimated value (greedy) and then the agent makes decision based on current information
```

▼ Epsilon-Greedy Policy

```
# The Epsilon-Greedy policy is a popular approach used in ad optimization, a common application of multi-arm bandit algorithms in the field of machine learning.
# In this context, the goal is to maximize the click-through rate (CTR) or any other relevant metric for the ads displayed to users.
```

```
# Here's a simple explained example of how the Epsilon-Greedy policy works in ad optimization:
```

```
# Let's say you are running an online advertising platform, and you have 5 different ad variations (arms) that you want to test.
# Each ad variation corresponds to a different design or message, and you want to determine which ad performs the best in terms of getting clicks.
```

```
# Initialization: At the beginning of the experiment, you have no information about which ad performs best.
# So, you set the initial value of epsilon,  $\epsilon$  (the exploration rate), to a high value, say 0.9.
# This means you will explore 90% of the time and exploit the best-performing ad only 10% of the time.
```

```
# Experimentation: You start displaying the 5 ad variations to different users in a random order.
# When a user visits your platform, you randomly select an ad to show them based on the current value of  $\epsilon$ .
```

```
# Exploration Phase: With a high value of  $\epsilon$ , you focus on exploring different ad variations.
# This means you show each ad almost equally to different users. During this phase, you gather data on how each ad performs in terms of clicks.
```

```
# Exploitation Phase: As the experiment progresses, you start to get an idea of which ad performs better based on the data collected so far.
# The Epsilon-Greedy policy will start to exploit the best-known ad more often.
```

```
# Updating Epsilon: Over time, you gradually reduce the value of  $\epsilon$ .
# For example, you may decay  $\epsilon$  by a small factor after a fixed number of rounds or based on the number of samples collected.
# As  $\epsilon$  decreases, the algorithm shifts from exploration to exploitation.
```

```
# Balancing Exploration and Exploitation: The key challenge is to find the right balance between exploration and exploitation.
# If  $\epsilon$  is too high for too long, you may not exploit the best-performing ad enough, resulting in suboptimal performance.
# On the other hand, if  $\epsilon$  decreases too quickly, you might miss out on finding potentially better-performing ads.
```

```
# Convergence: As the experiment continues and  $\epsilon$  approaches zero, the Epsilon-Greedy policy should converge towards exploiting the best-performing ad.
# This ad will be the one that maximizes the click-through rate or any other relevant metric.
```

```
# Continuous Learning: In a real-world scenario, you would continuously collect data and update your ad optimization strategy.
# New ads may be introduced, or ad performance may change over time, requiring the algorithm to adapt.
```

```
# By using the Epsilon-Greedy policy, you can efficiently explore different ad variations while exploiting the best-performing ad in the process,
# leading to an improved click-through rate and better ad optimization.
```

```
from io import IncrementalNewlineDecoder
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
np.random.seed(5)
```

```
n=10
arms=np.random.rand(n)
eps=0.1 # probability of exploration action
# eps is a very minimum value ( this is done to avoid the gradient descent problem)
# the epsilon value is minimal to strike a balance between exploring new options and exploiting the best-known option.
# When the epsilon value is minimal (close to zero), it means that the algorithm has mostly explored all the available options and
# has gathered enough information about their performance.

# At this point, the algorithm starts to shift its focus more towards exploitation, meaning it will choose the option that has shown the
# highest rewards or success.

# By minimizing the epsilon value, the algorithm becomes more confident in its decisions and starts to exploit the best option more frequently.
# This is beneficial because it allows the algorithm to take advantage of the knowledge it has gained through exploration and
# concentrate on making choices that are likely to yield the highest rewards or success.

# In summary, as the epsilon value becomes minimal, the algorithm becomes more selective in its actions, choosing the best-known option more often
# reducing its reliance on random exploration. This enables the algorithm to make more informed decisions and
# optimize its performance based on the information it has gathered during the exploration phase.
```

```
# this function will be used during the exploitation phase
def reward(prob): # probability of choosing an arm to exploit
```

```

reward=0
for i in range(10):
    if random.random() < prob: # this is a random value which is below the probability exploitation
        reward +=1
    return reward

# The function "reward" simulates the process of selecting an arm (option) to exploit during the exploitation phase.
# It does this by running a loop 10 times (you can change this number as needed).

# Inside the loop, it generates a random number between 0 and 1. If this random number is less than the probability "prob," it means the
# the agent gets a reward. The rewards obtained during the 10 iterations are added up, and the total rewards obtained during the exploitation phase.

# Remember, the higher the probability "prob," the more likely the arm is to be chosen for exploitation, and thus, the higher the expected rewards.
# The function simulates this process, allowing us to evaluate how well an arm performs in terms of rewards during the exploitation phase.

import random

# This function calculates the rewards obtained during the exploitation phase.
# It simulates selecting an arm (option) to exploit based on the given probability.

def reward(prob): # Define a function named "reward" that takes the probability "prob" as input.

    reward = 0 # Initialize a variable to keep track of the total rewards obtained.

    for i in range(10): # Repeat the following process 10 times (for example).

        if random.random() < prob: # Generate a random number between 0 and 1, and compare it with the probability "prob".
            # If the randomly generated number is less than the probability "prob"...
            # (In other words, there's a chance of "prob" that the arm will be selected for exploitation.)

            reward += 1 # Increment the rewards by 1, as the arm is chosen and the agent gets a reward.

    return reward # Return the total rewards obtained during the exploitation phase.

# initialize the memory array ; has 1 row defaulted to random activity index
av = np.array([random.randint(0,(n+1)),0]).reshape(1,2)

# random.randint(0, (n+1)): This generates a random integer between 0 and (n+1). Here, n represents the number of available actions or options.
# [random.randint(0, (n+1)), 0]: This creates a Python list with two elements. The first element is the randomly generated integer from 0 to n, representing the index of the action the agent chose. The second element is set to 0, representing the initial reward obtained for that action.
# np.array(...): This converts the Python list into a NumPy array, which is a more efficient data structure for numerical computations.
# reshape(1, 2): This reshapes the 1D array into a 2D array with one row and two columns. The first column represents the index of the action, and the second column represents the corresponding reward obtained for that action.

# Code initializes the av memory array with one row that stores the index of a randomly chosen action and its corresponding reward (which is 0).
# This memory array will be used to keep track of the agent's previous actions and rewards, enabling it to learn from its past experience.
# As the agent explores and exploits different actions, this memory will be updated to reflect the rewards received for each action.

# greedy method to select the best arm based on memory array
def bestArm(a):
    bestArm=0 # default 0
    bestMean=0
    for u in a:
        # calculate average reward for each action
        avg = np.mean(a[np.where(a[:,0]==u[0])][:,1])
        if bestMean < avg:
            bestMean = avg
            bestArm = u[0]
    return bestArm

# the function bestArm takes the memory array a as input and calculates the average reward for each action based on the history of rewards.
# It then identifies the action with the highest average reward, indicating the best-performing arm so far.

# By using the greedy approach, the function selects the arm with the highest average reward at the moment, without considering the potential for future rewards.
# The best arm is determined solely based on the historical data available in the memory array a.
# This method is simple and efficient but may not always lead to the globally optimal choice, especially in dynamic or non-stationary environments.
# Nevertheless, it can serve as a good starting point for exploring and exploiting actions in a multi-arm bandit problem.

def bestArm(a):
    bestArm = 0 # Initialize the index of the best arm to 0 (by default).
    bestMean = 0 # Initialize the mean reward of the best arm to 0.

    for u in a: # Loop through each row in the memory array `a`.
        # Calculate the average reward for each action based on the memory array.

```

```

# a[np.where(a[:, 0] == u[0])]: Select rows from the memory array where the first column (action index) matches the current action index
#[:, 1]: Select the second column (reward) from the rows that match the action index.

avg = np.mean(a[np.where(a[:, 0] == u[0])][:, 1])

# Compare the calculated average reward (`avg`) with the current best mean reward (`bestMean`).
# If the average reward for the current action is higher, update the best mean reward and the index of the best arm.
if bestMean < avg:
    bestMean = avg
    bestArm = u[0]

return bestArm # Return the index of the best arm (the action with the highest average reward).

# The next function you define is your greedy strategy of choosing the best arm so far. This function accepts a memory array that stores
# It is a 2 x k matrix where each row is an index reference to your arms array ( 1 st element), and the reward received ( 2nd element ).
# For example , if a row in your memory array is [2,8] it means that action 2 was taken ( the 3rd element in your arms array) and you received
# And here is the main loop for each play, let's play it 500 times and display a matplotlib scatter plot of the mean reward against the number of times played.

plt.xlabel("Number of times played")
plt.ylabel("Average Reward")

#Epsilon-greedy algorithm
for i in range(500):
    if random.random() > eps:
        #greedy exploitation action eps=0.1, 1-eps=0.9
        choice = bestArm(av) #call the method bestArm(action-value)
        thisAV = np.array([[choice, reward(arms[choice])]]) #call the reward method
        av = np.concatenate((av, thisAV), axis=0)
    else: #exploration - Agent explores other arms
        choice = np.where(arms == np.random.choice(arms))[0][0] #choice, reward
        thisAV = np.array([[choice, reward(arms[choice])]]) #add to our action value memory
    #calculate the mean reward
    runningMean = np.mean(av[:, 1])
    plt.scatter(i, runningMean)

```



